

数据隐私方法伦理和实践

Methodology, Ethics and Practice of Data Privacy

秘密共享, 安全多方计算

Secret Sharing, MPC

张兰

中国科学技术大学 计算机学院


2020春季

1. Secret Sharing

Threshold Scheme

» A (t,w) -threshold scheme

- Sharing key K among a set of w users
- Any t users can recover the key
- Any $t-1$ users can not do so



“

The security definition of secret sharing is that the shares must leak absolutely **no information** about the secret, until the number of shares passes the threshold value.

A Simple 2-out-of-2 Scheme

» Construction 3.5(2-out-of-2 TSSS)

$$\begin{array}{lll} \mathcal{M} = \{0, 1\}^\ell & \text{Share}(m): & \text{Reconstruct}(s_1, s_2): \\ t = 2 & s_1 \leftarrow \{0, 1\}^\ell & \text{return } s_1 \oplus s_2 \\ n = 2 & s_2 := s_1 \oplus m & \\ & \text{return } (s_1, s_2) & \end{array}$$

Example If we want to share the string $m = 1101010001$ then the Share algorithm might choose

$$s_1 := 0110000011$$

$$s_2 := s_1 \oplus m$$

$$= 0110000011 \oplus 1101010001 = 1011010010.$$

Then the secret can be reconstructed by xoring the two shares together, via:

$$s_1 \oplus s_2 = 0110000011 \oplus 1011010010 = 1101010001 = m.$$

Threshold Scheme

» A (t,w) -threshold scheme

- Sharing key K among a set of w users
- Any t users can recover the key
- Any $t-1$ users can not do so


» Schemes

- A Simple 2-out-of-2 Scheme
- Polynomial Interpolation
- Shamir Secret Sharing
- Visual Secret Sharing

Polynomial Interpolation

- » **Principle:** $d + 1$ points determine a unique degree- d polynomial.
- » If f is a polynomial that can be written as $f(x) = \sum_{i=0}^d f_i x^i$, then we say that f is a **degree- d** polynomial.

Polynomial Interpolation



Polynomials Over the Reals

» **Theorem 3.8 (Poly Interpolation)** Let

$\{(x_1, y_1), \dots, (x_{d+1}, y_{d+1})\} \subseteq \mathbb{R}^2$ be a set of points whose x_i values are all distinct. Then there is a unique degree- d polynomial f with real coefficients that satisfies $y_i = f(x_i)$ for all i .

Shamir Secret Sharing

- » **Idea.** We have just seen that any $d + 1$ points on a degree- d polynomial are enough to uniquely reconstruct the polynomial. So a natural approach for secret sharing is to let each user's share be a point on a polynomial.
- » **Shamir Secret Sharing** Share a secret $m \in Z_p$ with threshold t , first choose a degree- $(t - 1)$ polynomial f that satisfies $f(0) \equiv m \bmod p$, with all other coefficients chosen uniformly in Z_p . The i -th user receives the point $(i, f(i) \bmod p)$ on the polynomial.

Shamir Secret Sharing


» Example

Here is an example of 3-out-of-5 secret sharing over \mathbb{Z}_{11} (so $p = 11$). Suppose the secret being shared is $m = 7 \in \mathbb{Z}_{11}$. The Share algorithm chooses a random degree-2 polynomial with constant coefficient 7.

Let's say that the remaining two coefficients are chosen as $f_2 = 1$ and $f_1 = 4$, resulting in the following polynomial:

$$f(x) = 1x^2 + 4x + 7$$

This is the same polynomial illustrated in the previous example:



Shamir Secret Sharing

» Example(Continued)

For each user $i \in \{1, \dots, 5\}$, we distribute the share $(i, f(i) \% 11)$. These shares correspond to the highlighted points in the mod-11 picture above.


<i>user (i)</i>	<i>$f(i)$</i>	<i>share ($i, f(i) \% 11$)</i>
1	$f(1) = 12$	(1, 1)
2	$f(2) = 19$	(2, 8)
3	$f(3) = 28$	(3, 6)
4	$f(4) = 39$	(4, 6)
5	$f(5) = 52$	(5, 8)

- » Because Share is a randomized algorithm, there are **many valid** sharings of the same secret.

Blakley's Scheme

- » Secret is a point in an t -dimensional space
- » Dealer gives each user a hyper-plane passing the secret point
- » Any t users can recover the common point

Geometry View



Visual Secret Sharing

» a simple visual secret sharing scheme that is inspired by the following observations:

when  is stacked on top of , the result is 

when  is stacked on top of , the result is 

when  is stacked on top of , the result is 

when  is stacked on top of , the result is 

Visual Secret Sharing

» Construction 3.14

Share(m):

initialize empty images s_1, s_2 , with dimensions twice that of m

for each position (i, j) in m :

randomly choose $b_1 \leftarrow \{\text{white}, \text{black}\}$

if $m[i, j]$ is a white pixel: set $b_2 := b_1$

if $m[i, j]$ is a black pixel: set b_2 to the “opposite” of b_1 (i.e., $\{\text{white}, \text{black}\} \setminus \{b_1\}$)


add 2×2 block b_1 to image s_1 at position $(2i, 2j)$

add 2×2 block b_2 to image s_2 at position $(2i, 2j)$

return (s_1, s_2)


Visual Secret Sharing

» Example



Oblivious Transfer

- » Pick one out of two, without revealing which
 - Intuitive property: transfer partial information “obliviously”



2. MPC

Scenario 1: Private Dating

- » Alice and Bob meet at a pub
 - If both of them want to date together – they will find out
 - If Alice doesn't want to date – she won't learn his intentions
 - If Bob doesn't want to date – he won't learn her intentions



Scenario 1: Private Dating

- » Alice and Bob meet at a pub
- If both of them want to date together – they will find out
 - If Alice doesn't want to date – she won't learn his intentions
 - If Bob doesn't want to date – he won't learn her intentions

Solution: use a trusted bartender



Scenario 2: Private Auction

- » Many parties wish to execute a private auction
 - The highest bid wins
 - Only the highest bid (and bidder) is revealed



Scenario 2: Private Auction

- » Many parties wish to execute a private auction
 - The highest bid wins
 - Only the highest bid (and bidder) is revealed

Solution: use a trusted auctioneer



Scenario 3: Private Set Intersection

- » Intelligence agencies holds lists of potential terrorists
- The would like to compute the intersection
 - Any other information must remain secret



Mossad



MI5



FBI

Scenario 3: Private Set Intersection

- » Intelligence agencies holds lists of potential terrorists
- The would like to compute the intersection
 - Any other information must remain secret

Solution: use a trusted party



Mossad



MI5



FBI



Trust me

Scenario 4: Online Poker

- » Play online poker reliably



Scenario 4: Online Poker

» Play online poker reliably

Solution: use a trusted party




Secure Multiparty Computation

» **Goal:** use a protocol to emulate the trusted party

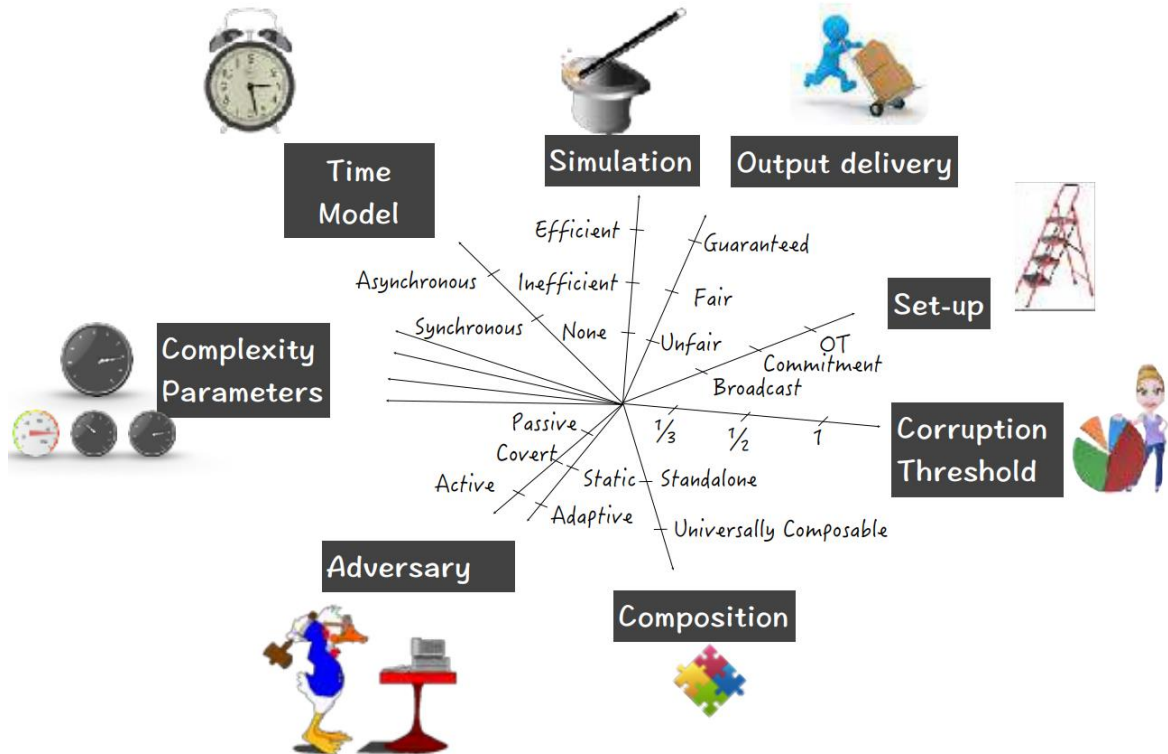


Secure Multiparty Computation

29




MPC dimensions




Ideal World

- » Each party sends its input to the trusted party
- » The trusted party computes $y=f(x_1,\dots,x_n)$
- » Trusted party sends y to each party




Real World

- » Parties run a protocol π on inputs (x_1, \dots, x_n)

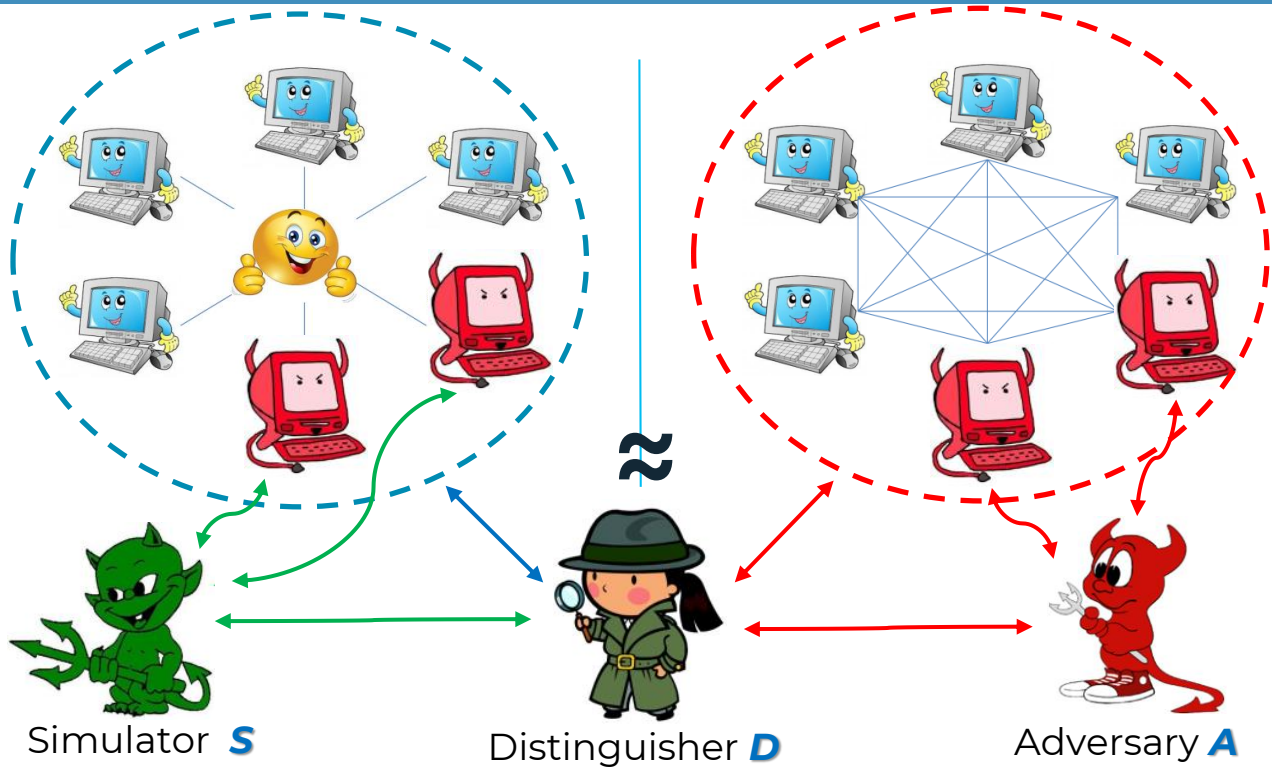


Simulation-Based Security



Distinguisher **D**

Simulation-Based Security



Simulation-Based Security

- » The distinguisher D :
 - Gives inputs to parties
 - Gets back output from parties and from adversary/simulator
 - Guesses which world it is real/ideal
- » Protocol π securely computes y if $\forall A \exists S \forall D$ distinguishing success is “small”



For every real-model adversary A , there exists an ideal-model simulator S , such that the result of a real execution of π with A is indistinguishable from the result of an ideal execution with S

The Definition

- » A definition of an MPC task involves defining:
 - **Functionality:** what do we want to compute?
 - **Security type:** how strong protection do we want?
 - **Adversarial model:** what do we want to protect against?
 - **Network model:** in what setting are we going to do it?

The Functionality

- » The code of the trusted party
- » We will focus on secure function evaluation (SFE), the trusted party computes $y=f(x_1,\dots,x_n)$
 - Deterministic vs. randomized
 - Single public output vs. private outputs
 - Reactive vs. non-reactive

Security Type

» **Computational:**

- The real & ideal worlds are computationally indistinguishable

» **Statistical (information-theoretic):** all-powerful distinguisher, negligible error probability

- The real & ideal worlds are statistically close

» **Perfect:** all-powerful distinguisher, zero error probability

- The real & ideal worlds are identically distributed

Adversarial Model (1)

» Adversarial behavior

- **Semi honest:** honest-but-curious. corrupted parties follow the protocol honestly, adversary tries to learn more information.
- **Fail stop:** same as semi honest, but corrupted parties can prematurely halt.
- **Malicious:** corrupted parties can deviate from the protocol in an arbitrary way

Adversarial Model (2)

» Adversarial power

- **Polynomial time:** computational security, normally requires cryptographic assumptions, e.g., encryption, signatures, oblivious transfer
- **Computationally unbounded:** an all-powerful adversary, information-theoretic security

Adversarial Model (3)

» Adversarial corruption

- **Static:** the set of corrupted parties is defined before the execution of the protocol begins. Honest parties are always honest, corrupted parties are always corrupted
- **Adaptive:** adversary can decide which parties to corrupt during the course of the protocol, based on information it dynamically learns
- **Mobile:** adversary can “jump” between parties. Honest parties can become corrupted, corrupted parties can become honest again

Adversarial Model (4)

- » Number of corrupted parties
 - **Threshold adversary:** Denote by $t \leq n$, an upper bound on # corruptions
 - » No honest majority, e.g., two-party computation
 - » Honest majority, i.e., $t \leq n/2$
 - » Two-thirds majority, i.e., $t \leq n/3$
 - **General adversary structure:**
 - » Protection against specific subsets of parties

Communication Model (1)

- » **Point-to-point:** fully connected network of pairwise channels. (Partial networks: star, chain)
 - Unauthenticated channels
 - **Authenticated channels:** in the computational setting
 - **Private channels:** in the IT setting
- » **Broadcast:** additional broadcast channel

Communication Model (2)

» Message delivery:

- **Synchronous:** the protocol proceeds in rounds. Every message that is sent arrives within a known time frame
- **Asynchronous (eventual delivery):** the adversary can impose arbitrary (finite) delay on any message
- **Fully Asynchronous:** the adversary has full control over the network, can even drop messages

Execution Environment

» Stand alone:

- A single protocol execution at any given time (isolated from the rest of the world)

» Concurrent general composition:

- Arbitrary protocols are executed concurrently
- An Internet-like setting
- Requires a strictly stronger definition. Captured by the universal composability (UC) framework
- Impossible in general without a trusted setup assumption (e.g., common reference string)

Feasibility Results

» Malicious setting

- For $t \leq n/3$, every f can be securely computed with **perfect** security [BGW'88, CCD'88]
- For $t \leq n/2$, every f can be securely computed with **statistical** security [RB'89]
- For $t \leq n$, assuming OT, every f can be securely computed with **abort and computational** [GMW'87]

» Semi-honest setting

- For $t \leq n/2$, every f can be securely computed with **perfect** security [BGW'88, CCD'88]
- For $t \leq n$, assuming OT, every f can be securely computed with **computational** security [GMW'87]
-

Defining Privacy

- » First pick an input (x_1, \dots, x_n) for the protocol and make a run of the protocol on this input.
- » Then pick some $C \subset \{P_1, \dots, P_n\}$ with $|C| \leq t$ and consider the values $\{view_j\}_{P_j \in C}$, where $view_j$ is the view of party P_j in the execution.
- » **leaked** values : the values $\{view_j\}_{P_j \in C}$.
- » **allowed** values: the values $\{x_j, y_j\}_{P_j \in C}$, (the input and output of each party)

Defining Privacy

- » A protocol is **private** if it always holds that the leaked values contain no more information than the allowed values.
- » A protocol is private if it always holds that the leaked values can be computed efficiently from the allowed values.

MPC Techniques


- » Secure Multiparty Computation
 - Yao's Garbled Circuit
 - Secret sharing schemes
 - Homomorphic Encryption

Yao's Theorem

- » The first completeness theorem for secure computation.
- » It states that for **ANY** efficiently computable function, there is a secure two-party protocol in the semi-honest model.
 - Therefore, in theory there is no need to design protocols for specific functions.
 - Surprising!


Yao's Protocol

- » Compute **any** function securely
 - ... in the semi-honest model
- » First, convert the function into a **boolean circuit**




1: Pick Random Keys For Each Wire

- » Next, evaluate one gate securely
 - Later, generalize to the entire circuit
- » Alice picks two **random keys** for each wire
 - One key corresponds to “0”, the other to “1”
 - 6 keys in total for a gate with 2 input wires



2: Encrypt Truth Table

- » Alice encrypts each row of the truth table by encrypting the output-wire key with the corresponding pair of input-wire keys



Original truth table:

x	y	z
0	0	0
0	1	0
1	0	0
1	1	1


Encrypted truth table:

$$\begin{aligned} &E_{k_{0x}}(E_{k_{0y}}(k_{0z})) \\ &E_{k_{0x}}(E_{k_{1y}}(k_{0z})) \\ &E_{k_{1x}}(E_{k_{0y}}(k_{0z})) \\ &E_{k_{1x}}(E_{k_{1y}}(k_{1z})) \end{aligned}$$

3: Send Garbled Truth Table

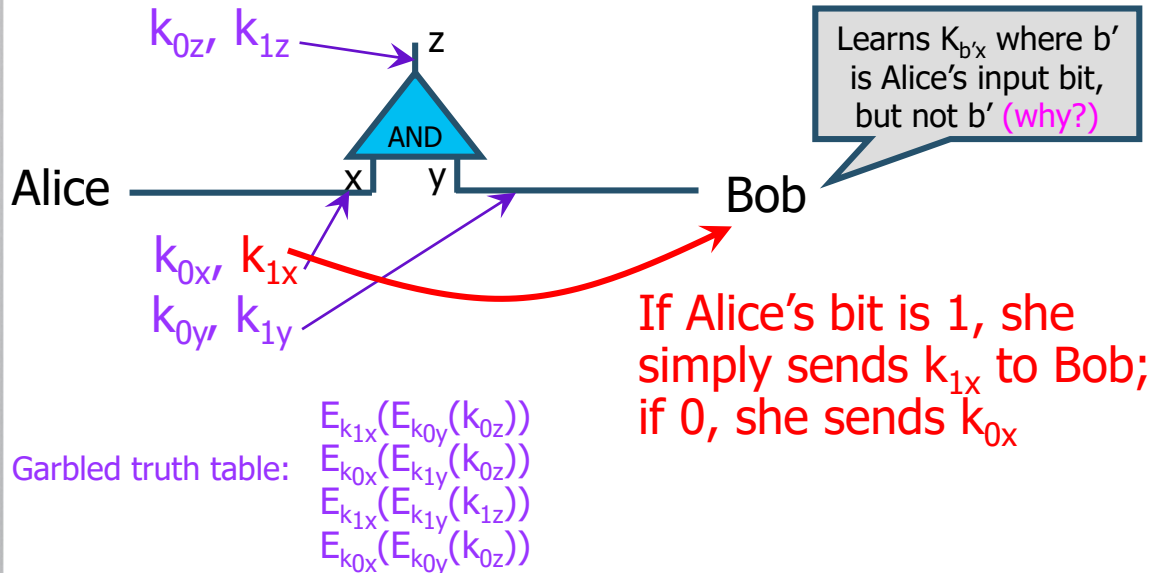
- » Alice randomly permutes (“garbles”) encrypted truth table and sends it to Bob

Does not know which row of garbled table corresponds to which row of original table



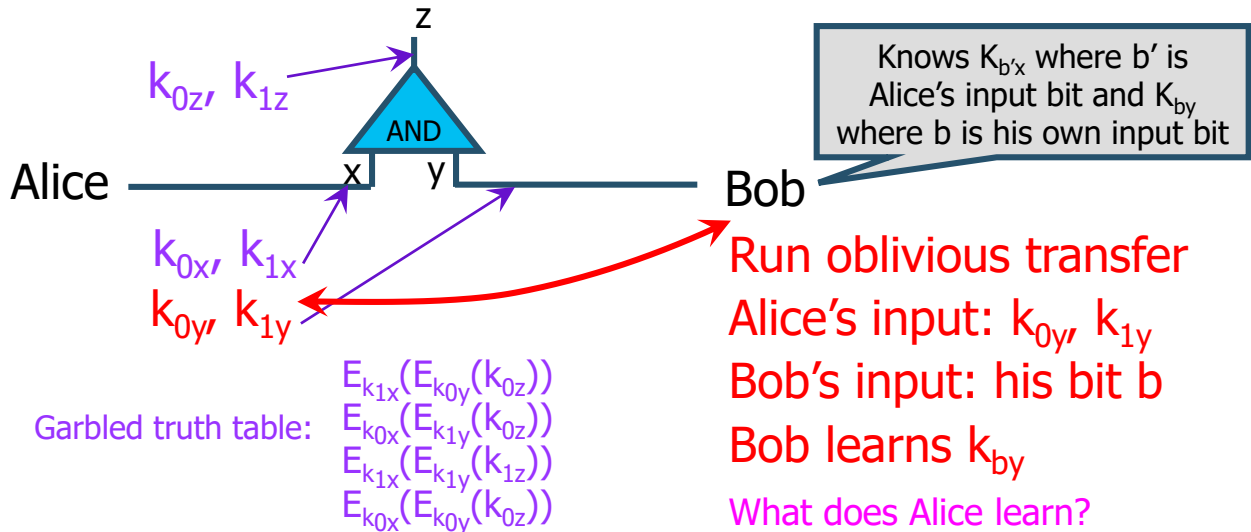
4: Send Keys For Alice's Inputs

- » Alice sends the key corresponding to her input bit
 - Keys are random, so Bob does not learn what this bit is



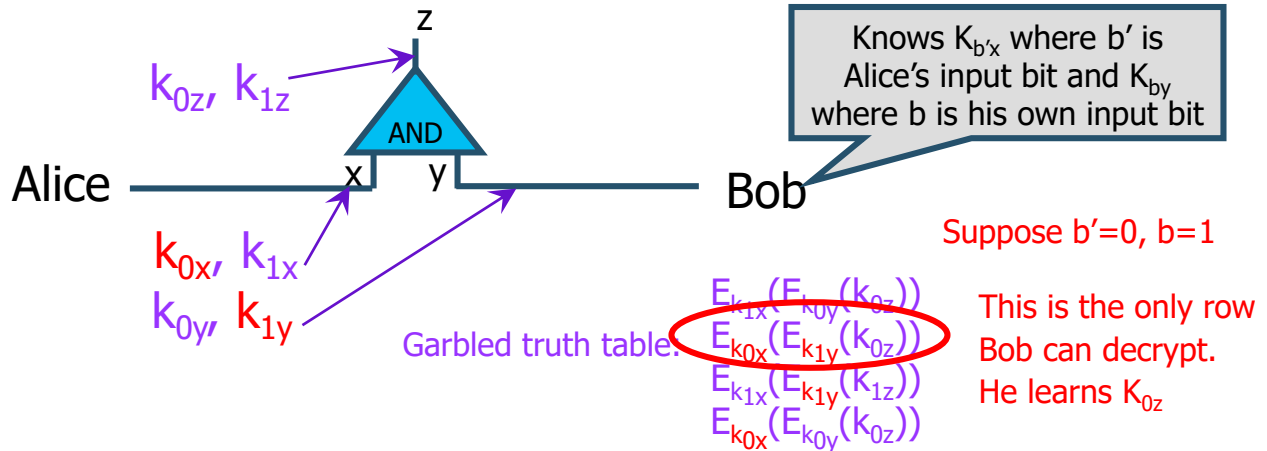
5: Use OT on Keys for Bob's Input

- » Alice and Bob run oblivious transfer protocol
- Alice's input is the two keys corresponding to Bob's wire
 - Bob's input into OT is simply his 1-bit input on that wire



6: Evaluate Garbled Gate

- » Using the two keys that he learned, Bob decrypts exactly one of the output-wire keys
 - Bob does not learn if this key corresponds to 0 or 1
 - » Why is this important?




Evaluation of Garbed Circuit

- » Given the two keys (one for Alice and one for Bob) representing the inputs of a gate, we can easily obtain the key representing the output of the gate.
 - Only need to decrypt the corresponding entry using both keys from Alice and Bob
 - But we do not know which entry it is? We can decrypt all entries. Suppose each cleartext contains some redundancy (like a hash value). Then only decryption of the right entry can yield such redundancy.

7: Evaluate Entire Circuit

- » In this way, Bob evaluates entire garbled circuit
 - For each wire in the circuit, Bob learns only one key
 - It corresponds to 0 or 1 (Bob does not know which)
 - » Therefore, Bob does not learn intermediate values (why?)
- » Bob tells Alice the key for the final output wire and she tells him if it corresponds to 0 or 1
 - Bob does not tell her intermediate wire keys (why?)



Finishing the Evaluation

- » At the end of evaluation, Bob gets the keys representing the output bits of circuit.
 - Alice sends Bob a table of the keys for each output bit.
 - Bob translates the keys back to the output bits.
- » For privacy, we need to be careful:
 - The topology of the circuit should be the same for all circuits of a particular input size.
 - Then privacy is guaranteed.

Brief Discussion of Yao's Protocol

- » Function must be converted into a circuit
 - For many functions, circuit will be huge
- » If m gates in the circuit and n inputs, then need $4m$ encryptions and n oblivious transfers
 - Oblivious transfers for all inputs can be done in parallel
- » Yao's construction gives a constant-round protocol for secure computation of any function in the semi-honest model
 - Number of rounds does not depend on the number of inputs or the size of the circuit!

How to ensure Efficiency?

» Computation and Communication?

- Communication optimization: reducing information that must be shared between the two parties;
- Execution optimization: allowing for the same number of gates to be executed in a shorter amount of time.
- Circuit optimization: reducing the number of gates needed to compute a function.

Privacy \neq Security

- » As we have mentioned, secure computation only deals with the process of computing the function
 - It does not ask whether or not the function should be computed

Privacy and Secure Computation

- » Secure computation can be used to solve any distributed data-mining problem
- » A two-stage process:
 - Decide that the function/algorithm should be computed – an issue of **privacy**
 - Apply secure computation techniques to compute it securely – **security**
- » But, not every privacy problem can be cast as a distributed computation

THANKS!

Any questions?

You can find me at:

» zhanglan@ustc.edu.cn

